

Managing Terraform State

The Plan

- 🚚 What is Terraform state?
- 🚚 Shared storage for state files
- 🚚 Limitations with Terraform's backends
- 🚚 Isolating state files
 - Isolation via workspaces
 - Isolation via file layout
- 🚚 The `terraform_remote_state` data source

What Is Terraform State?

- 🚚 Every time you run Terraform, it records information about what infrastructure it created in a Terraform state file
- 🚚 By default, when you run Terraform in the folder /foo/bar, Terraform creates the file /foo/bar/terraform.tfstate.
- 🚚 Say, your Terraform configuration shows below

```
1 resource "aws_instance" "example" {  
2   ami           = "ami-0c55b159cbfaffe1f0"  
3   instance_type = "t2.micro"  
4 }
```

Example of "terraform.tfstate"



After you run `terraform apply` you will see the output in

```
t {
  "version": 4,
  "terraform_version": "0.12.0",
  "serial": 1,
  "lineage": "1f2087f9-4b3c-1b66-65db-8b78faafc6fb",
  "outputs": {},
  "resources": [
    {
      "mode": "managed",
      "type": "aws_instance",
      "name": "example",
      "provider": "provider.aws",
      "instances": [
        {
          "schema_version": 1,
          "attributes": {
            "ami": "ami-0c55b159cbfafa1f0",
            "availability_zone": "us-east-2c",
            "id": "i-00d689a0acc43af0f",
            "instance_state": "running",
            "instance_type": "t2.micro",
            "(...)": "(truncated)"
          }
        }
      ]
    }
  ]
}
```

Meaning of "terraform.tfstate"

- 🚚 Resource with `type aws_instance` and `name example` corresponds to an EC2 Instance in your AWS account with ID `i-00d689a0acc43af0f`
- 🚚 Every time you run Terraform
 - it can fetch the latest status of this EC2 Instance from AWS
 - compare that to what's in your Terraform configurations
 - determine what changes need to be applied
- 🚚 Thus, the output of the `terraform plan` command is a diff
 - between the code on your computer and
 - the infrastructure deployed in the real world, as discovered via IDs in the state file

Managing Terraform States

- 🚚 State file tracks managed resources created by Terraform
- 🚚 Other resources are ignored unless there is conflict
- 🚚 Existing resources can be added to a Terraform state
- 🚚 Resources can be removed from a Terraform state
- 🚚 Multiple state files can be used
- 🚚 Each state file manages a only its set of resources

The "state" Command



The state command has multiple options (not all are listed)


- `terraform state list` : lists the resources being managed
- `terraform state show <resource>` : displays state data for a resource
- `terraform state rm <resource>` : stops managing the AWS object linked to <resource>



`terraform import <resource> <AWS ID>` : links the Terraform resource with a terrafrom resource

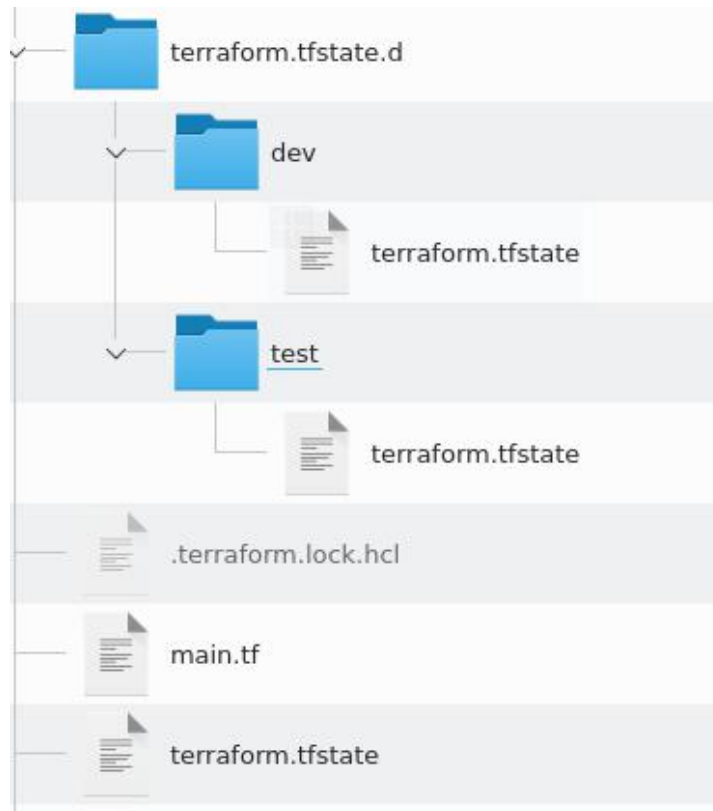
Separate Environments

- 🚚 We often need multiple copies of a deployment for different purposes
- 🚚 Common environments are: development, test, stage and production



Terraform Workspaces

- 📘 Terraform supports a separate configuration for each deployment
 - Each deployment is called a workspace
 - There is always a `default` workspace
- 📘 We can create additional workspaces as we need them
 - For example, we could have defined `dev` and `test` workspaces
- 📘 For local state files, each new workspace's state file is in its own folder



The "workspace" Command



`terraform workspace` has several options:

- `list` : lists all workspaces marking the current one with `"*`
- `show` : lists the currently active workspace
- `new <name>` : creates and switches to the newly created workspace
- `select <name>` : switches to the named workspace
- `delete <name>` : deletes the named workspace
 - The "default" workspace can never be deleted
 - Deleting a workspace does **not** destroy the resources, it just leaves them unmanaged

Remote Backends



Each Terraform configuration specifies where the state files are kept

- This is called the "backend"
- The default is to use files in the local directory
- This is what we have been using so far



Terraform can also support "remote" backends

- For example, we can keep state files in an S3 bucket on AWS
- Not all providers can host remote back ends

Problem with Local Backends



Shared storage for state files

- Files need to be in common shared area so everyone on the team can access them
- Without file locking, race conditions when concurrent updates to the state files take place
- This can lead to conflicts, data loss, and state file corruption



Isolation

- It's difficult to isolate the code used in different environments
- Lack of isolation makes it easy to accidentally overwrite environments



Secrets

- Confidential information is stored in the clear (i.e. AWS Keys)

Remote AWS Backend



Using S3 as a backend resolves many of these issues

- S3 manages the updating and access independently, and supports versions
- S3 supports encryption
- S3 supports locking for multiple access
- S3 allows a common repository we can control access to



S3 is also managed so that we don't have to manage it

- S3 has high levels of availability and durability
- S3 also means we have reduced the risk of "loosing" configurations

Setting up the S3 Bucket

```
1 resource "aws_s3_bucket" "terraform_state" {
2     bucket = "terraform-up-and-running-state"
3
4     # Prevent accidental deletion of this S3 bucket
5     lifecycle {
6         prevent_destroy = true
7     }
8
9     # Enable versioning so we can see the full revision history of our
10    # state files
11    versioning {
12        enabled = true
13    }
14
15    # Enable server-side encryption by default
16    server_side_encryption_configuration {
17        rule {
18            apply_server_side_encryption_by_default {
19                sse_algorithm = "AES256"
20            }
21        }
22    }
23 }
```

Setting up the Locking Table



Next, a DynamoDB table to use for locking

- DynamoDB is Amazon's distributed key-value store
- It supports strongly consistent reads and conditional writes

```
1 resource "aws_dynamodb_table" "terraform_locks" {  
2     name           = "terraform-up-and-running-locks"  
3     billing_mode    = "PAY_PER_REQUEST"  
4     hash_key        = "LockID"  
5  
6     attribute {  
7         name = "LockID"  
8         type  = "S"  
9     }  
10 }
```

Setting Up the Backend



We have to tell Terraform the backend is now remote

– We do this in the `terraform` directive

```
1 terraform {  
2   backend "< BACKEND_NAME >" {  
3     [CONFIG...]  
4   }  
5 }
```

```
1 terraform {  
2   backend "s3" {  
3     # Replace this with your bucket name!  
4     bucket      = "terraform-up-and-running-state"  
5     key         = "global/s3/terraform.tfstate"  
6     region      = "us-east-2"  
7  
8     # Replace this with your DynamoDB table name!  
9     dynamodb_table = "terraform-up-and-running-locks"  
10    encrypt      = true  
11  }  
12 }
```


Moving State File Locations



To move local state to a remote backend

- Create the remote backend resources and define the backend configuration
- Run `terraform init` and the local config is copied to the remote backend



To move from remote backend to a local backend

- Remove the backend configuration
- Run `terraform init` and the remote config is copied to the local backend

Moving Backends Summary



To make this work, you need to use a two-step process:

- Create the S3 bucket and DynamoDB table and deploy that code with a local backend
- Add a remote backend configuration to it to use the S3 bucket and DynamoDB table
- Run terraform init to copy your local state to S3



To revert to a local state backend

- Remove the backend configuration
- Rerun terraform init to copy the Terraform state to the local disk
- Run terraform destroy to delete the S3 bucket and DynamoDB table

Remote Backend Advantage

- 📦 A single S3 bucket and DynamoDB table can be shared across all your Terraform code
- 📦 You'll probably only need to do it once per AWS account
- 📦 After the S3 bucket exists, in the rest of your Terraform code, you can specify the backend configuration right from the start without any extra steps

Backend Limitation

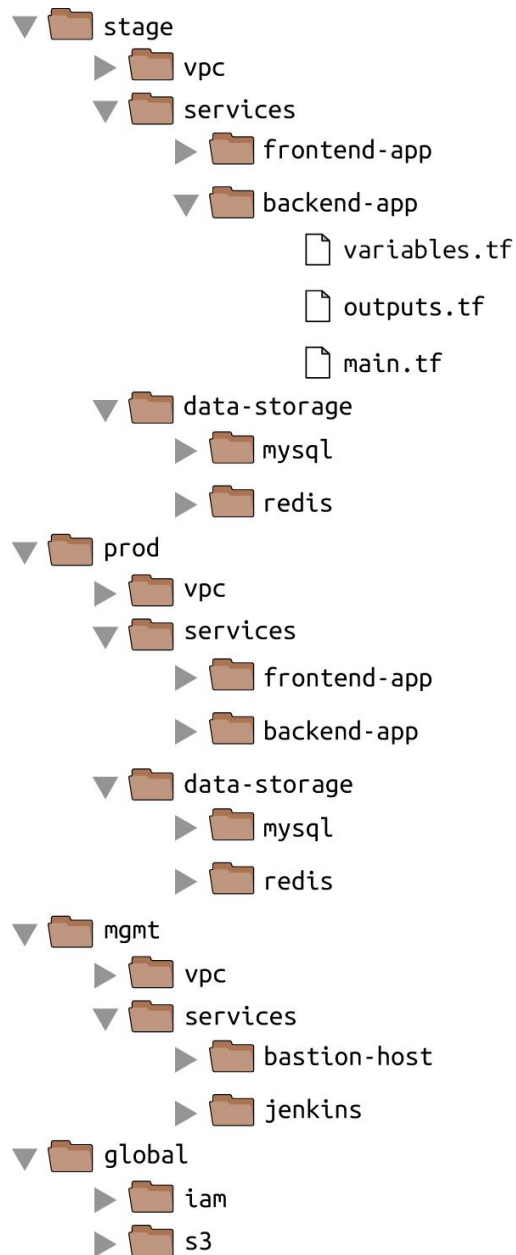
- 🚫 Variables and references cannot be used in the backend block
- 🚫 The following will **not** work

```
1 # This will NOT work. Variables aren't allowed in a backend configuration.
2 terraform {
3     backend "s3" {
4         bucket      = var.bucket
5         region      = var.region
6         dynamodb_table = var.dynamodb_table
7         key          = "example/terraform.tfstate"
8         encrypt      = true
9     }
10 }
```

File Isolation

- 🔒 Most secure approach is to have a folder for each configuration
- 🔒 Each deployment has its own backend, local or remote.
 - This allows for isolation of all files
 - Allows for different access and authentication mechanisms
 - Eg. Different S3 buckets used as backends can have different policies

File Isolation Example



Workspaces Use Case



If you already have a Terraform module deployed

- you want to do some experiments with it
- but you don't want your experiments to affect the state of the already deployed infrastructure



Run `terraform workspace new` to deploy a new copy of the exact same infrastructure, but storing the state in a separate file

Workspace Specific Configurations

- 🚀 You can even change how that module behaves based on the workspace you're in by reading the workspace name using the expression `terraform.workspace`

```
1 resource "aws_instance" "example" {  
2     ami          = "ami-0c55b159cbfafa1f0"  
3     instance_type = terraform.workspace == "default" ? "t2.medium" : "t2.micro"  
4 }
```

- 🚀 Workspaces allow a fast and easy way to quickly spin up and tear down different versions of your code

Workspace Drawbacks



All workspace state files are stored in the same backend

- They share same authentication and access controls which means they are not good for isolating



Workspaces are not visible in the code or on the terminal unless you run terraform workspace commands

- A module in one workspace looks exactly the same as a module deployed in 10 workspaces
- This makes maintenance more difficult, because you don't have a good picture of your infrastructure



Workspaces can be fairly error prone

- The lack of visibility makes it easy to forget what workspace you're in and accidentally make changes in the wrong one

Isolation via File Layout



To achieve full isolation between environments:

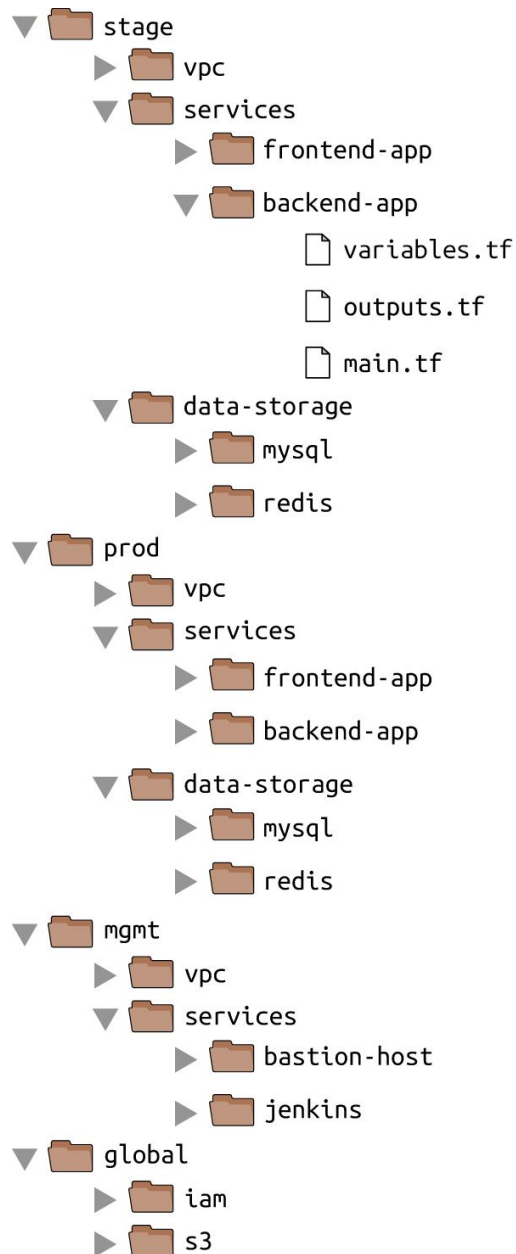
- Put the Terraform configuration files for each environment into a separate folder
- For example, all of the configurations for the staging environment can be in a folder called stage
- All the configurations for the production environment can be in a folder called prod



Configure a different backend for each environment, using different authentication mechanisms and access controls

- Each environment could live in a separate AWS account with a separate S3 bucket as a backend

Typical Project File Layout



Isolation via File Layout



At the top level, there are separate folders for each “environment

- **stage** : An environment for preproduction workloads (testing)
- **prod** : An environment for production workloads (user facing apps)
- **mgmt** : An environment for DevOps tooling (Jenkins etc.)
- **global** : Resources that are used across all environments (S3, IAM)



Within each environment, there are separate folders for each “component”:

- **vpc** : Network topology for this environment
- **services** : Apps or microservices to run in this environment - each app could have its own folder to isolate it
- **data-storage** : The data stores to run in this environment, such as MySQL or Redis

Isolation via File Layout



Within each component are the actual Terraform configuration files with the following naming conventions:

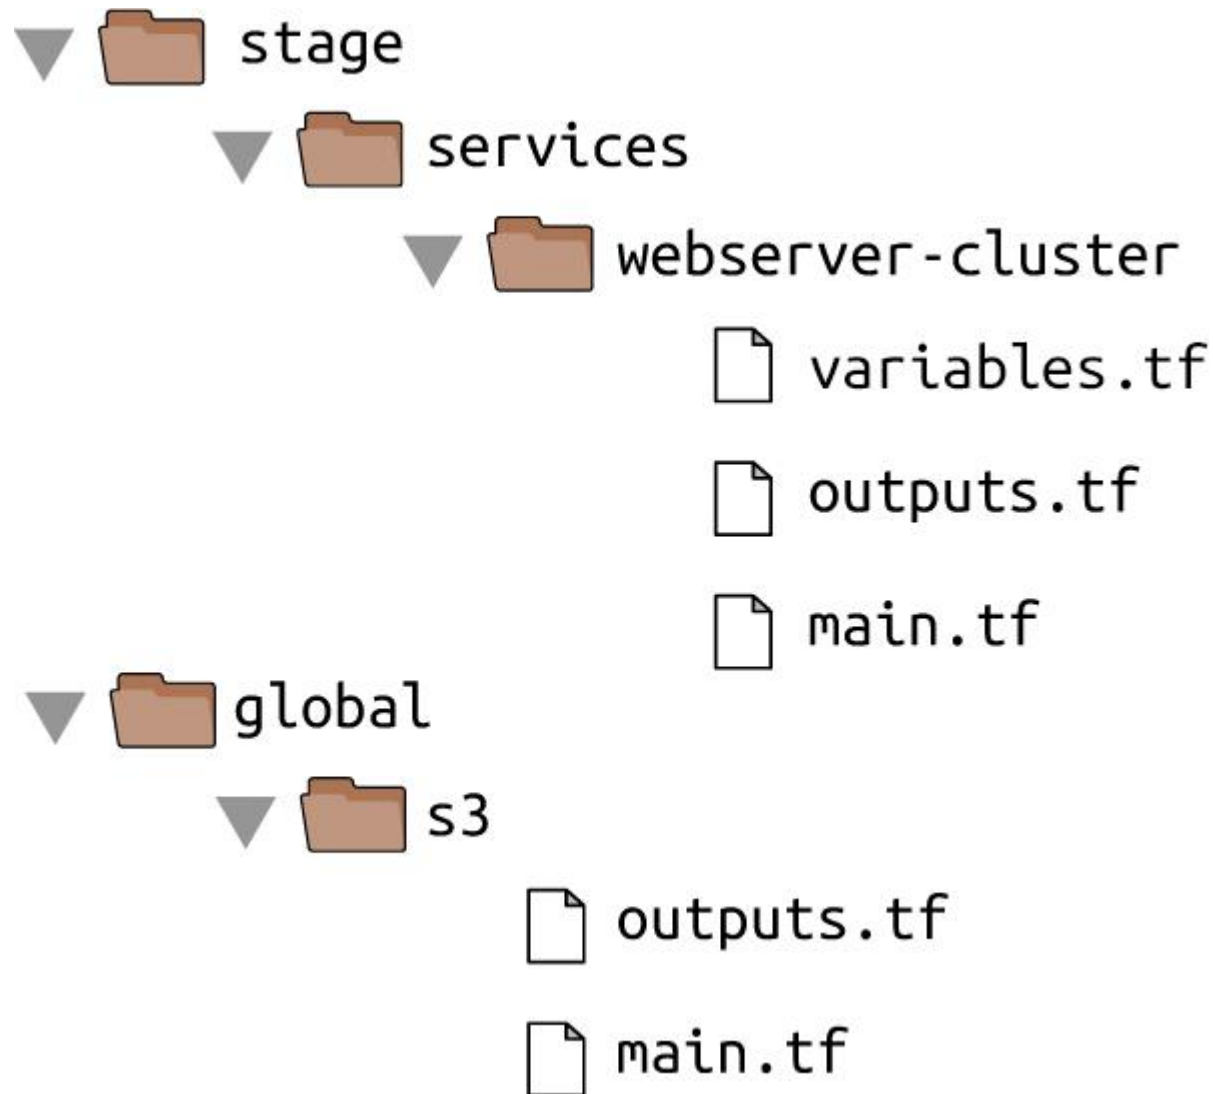
- **variables.tf** : Input variables
- **outputs.tf** : Output variables
- **main.tf** : The resources



Terraform looks for files in the current directory with the .tf extension

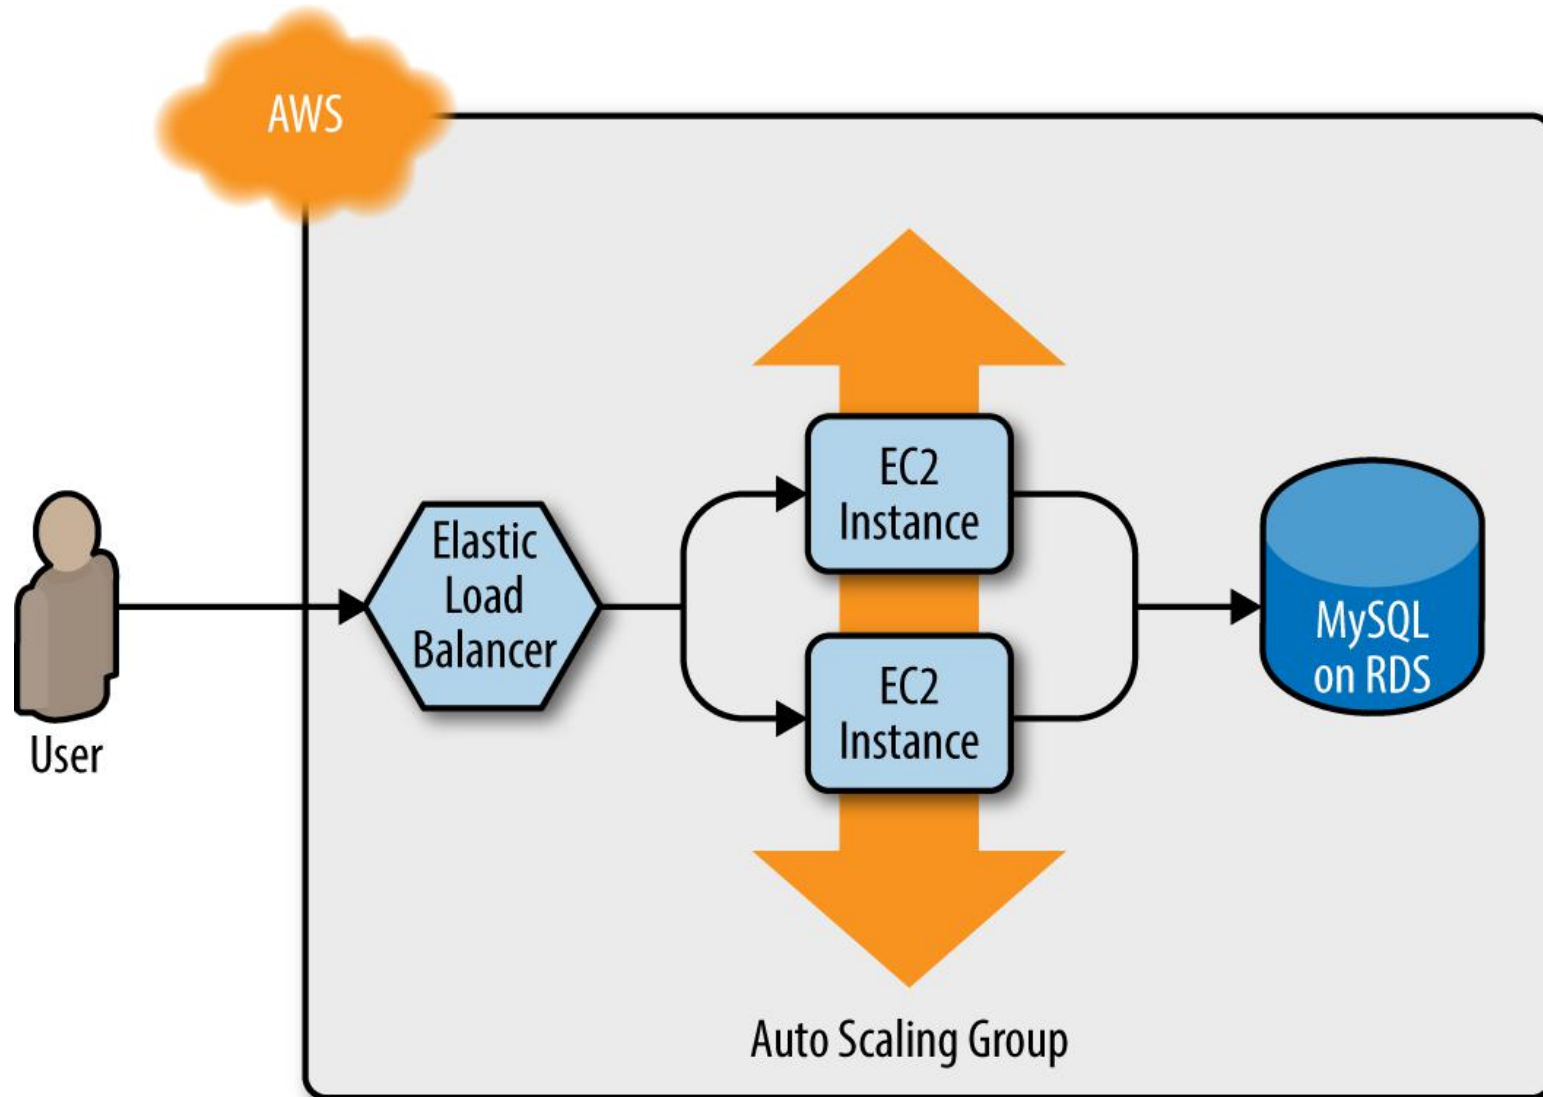
- Using a consistent, predictable naming convention makes code easier to browse
- Then you always know where to look to find a variable, output, or resource. If individual Terraform files are

Rarranged Sample Code



The "terraform_remote_state" Data Source

- Assume that the web server cluster needs to communicate with a MySQL database



Deployment Consideration

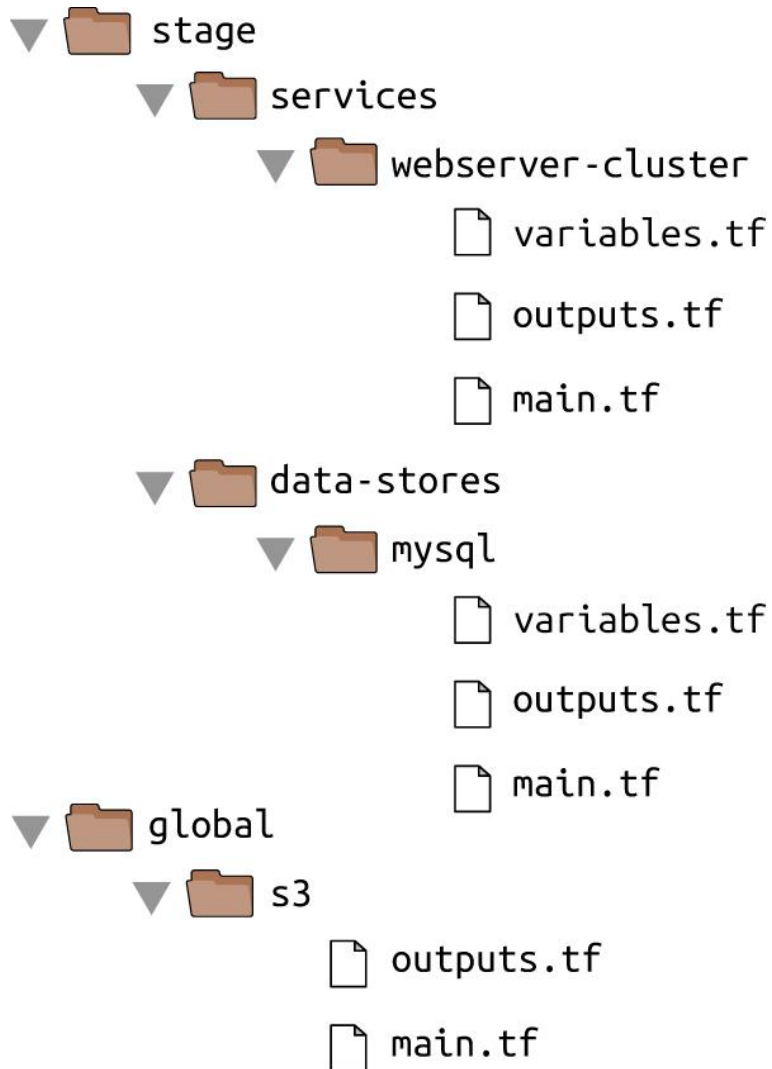


- The MySQL database should probably be managed with a different set of configuration files as the web server cluster
- Updates will probably be deployed to the web server cluster frequently
 - What to avoid accidentally breaking the database when doing an update

Deployment Consideration



Isolate the MySQL configurations in a data-stores folder



Keeping Secrets



One of the parameters that you must pass to the `aws_db_instance` resource is the master password to use for the database

- This should not be in the code in plain text
- There are two other options



Read the secret from a secret store - there are multiple secrets managers

- AWS Secrets Manager and the `aws_secretsmanager_secret_version` data source (shown in the example code)
- AWS Systems Manager Parameter Store and the `aws_ssm_parameter` data source
- AWS Key Management Service (AWS KMS) and the `aws_kms_secrets` data source
- Google Cloud KMS and the `google_kms_secret` data source
- Azure Key Vault and the `azurerm_key_vault_secret` data source
- HashiCorp Vault and the `vault_generic_secret` data source

Using AWS Secrets Manager



```
1 resource "aws_db_instance" "example" {
2     identifier_prefix    = "terraform-up-and-running"
3     engine               = "mysql"
4     allocated_storage    = 10
5     instance_class       = "db.t2.micro"
6     name                 = "example_database"
7     username             = "admin"
8
9     password =
10     data.aws_secretsmanager_secret_version.db_password.secret_string
11 }
12
13 data "aws_secretsmanager_secret_version" "db_password" {
14     secret_id = "mysql-master-password-stage"
15 }
```

Keeping Secrets II

🔒 Other option is to manage them completely outside of Terraform

- Then pass the secret into Terraform via an environment variable.
- In the code below, there is no default since it's a secret

```
1 variable "db_password" {  
2     description = "The password for the database"  
3     type        = string  
4 }  
5  
6 export TF_VAR_db_password="(YOUR_DB_PASSWORD)"  
7 $ terraform apply
```

🔒 A known weakness of Terraform:

- The secret will be stored in the Terraform state file in plain text
- The only solution is to lock down and encrypt the state files

Creating the Database



Using the backend to create the state repository:

```
1 terraform {
2   backend "s3" {
3     # Replace this with your bucket name!
4     bucket      = "terraform-up-and-running-state"
5     key         = "stage/data-stores/mysql/terraform.tfstate"
6     region      = "us-east-2"
7
8     # Replace this with your DynamoDB table name!
9     dynamodb_table = "terraform-up-and-running-locks"
10    encrypt       = true
11  }
12 }
```




Run the Terraform `init` and `apply` commands to create the database

- Provide the database ports to the webserver cluster

```
1 output "address" {
2   value      = aws_db_instance.example.address
3   description = "Connect to the database at this endpoint"
4 }
5
6 output "port" {
7   value      = aws_db_instance.example.port
8   description = "The port the database is listening on"
9 }
```

Integrating the Database

 Running `init` again produces:

```
1 $ terraform apply
2
3 (...)
4
5 Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
6
7 Outputs:
8
9 address = tf-2016111123.cowu6mts6srx.us-east-2.rds.amazonaws.com
10 port = 3306
```

 These outputs are now stored in the Terraform state for the database

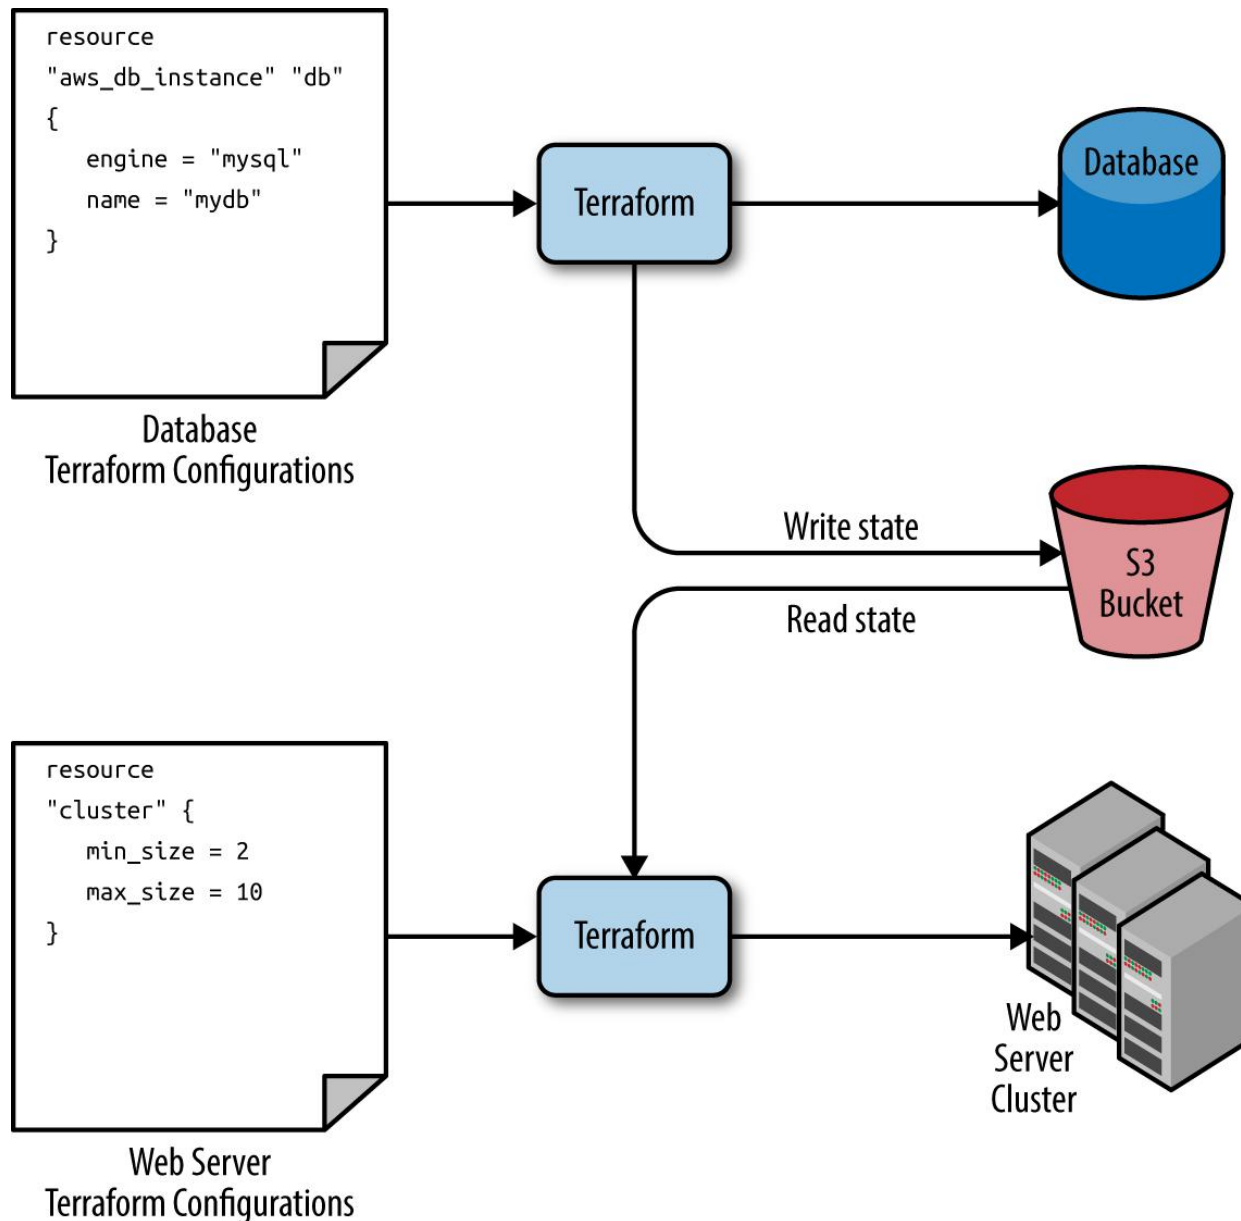
- The web server cluster code can read the data from this state file by adding the `terraform_remote_state` data source in `stage/services/webserver-cluster/main.tf`:

```
1 data "terraform_remote_state" "db" {
2     backend = "s3"
3
4     config = {
5         bucket = "(YOUR_BUCKET_NAME)"
6         key    = "stage7/data-stores/mysql/terraform.tfstate"
7         region = "us-east-2"
8     }
```

Reading the DB State



The web service cluster reads the DB configuration from the state file in the S3 bucket



Reading the DB State

- 📘 All Terraform data sources, like the data returned by `terraform_remote_state`, is read-only
 - Nothing in the webserver code can modify the state
 - The database's state data can be read with no risk to the database
- 📘 The output variables are stored in the state file can be read using an attribute reference of the form:

```
1 data.terraform_remote_state.< NAME >.outputs.< ATTRIBUTE >
```

- 📘 This code updates the User Data of the web server cluster Instances to pull the database address and port out of the `terraform_remote_state` data source:

```
1 user_data = << EOF # space added for formatting
2 #!/bin/bash
3 echo "Hello, World" >> index.html
4 echo "${data.terraform_remote_state.db.outputs.address}" >> index.html
5 echo "${data.terraform_remote_state.db.outputs.port}" >> index.html
6 nohup busybox httpd -f -p ${var.server_port} &
7 EOF
```


Template Files

- 🚗 Hard-coding the script in the previous slide is not effective
 - Instead we can read the contents from a file using the Terraform `file()` function which reads the contents of a file and returns it as a string
 - However, we have to dynamically insert Terraform data
- 🚗 To provide this facility, Terraform has a `template_file` data source that has two arguments: A template, which is a string to render and a map of variables to use while rendering as well as one output attribute called `rendered`, which is the result of rendering template
 - For example, this can be added to the `webserver-cluster`

```
1 data "template_file" "user_data" {
2     template = file("user-data.sh")
3
4     vars = {
5         server_port = var.server_port
6         db_address  = data.terraform_remote_state.db.outputs.address
7         db_port     = data.terraform_remote_state.db.outputs.port
8     }
9 }
```

Template Files



We also have to provide the "slots" in the template code for insertion of the variables

- We use standard Terraform string interpolation, and we don't need the `var` prefix

```
1 #!/bin/bash
2
3 cat > index.html << EOF # space added for formatting
4 <h1>Hello, World</h1>
5 <p>DB address: ${db_address}</p>
6 <p>DB port: ${db_port}</p>
7 EOF
8
9 nohup busybox httpd -f -p ${server_port} &
```

```
1 resource "aws_launch_configuration" "example" {
2     image_id      = "ami-0c55b159cbfaffe1f0"
3     instance_type = "t2.micro"
4     security_groups = [aws_security_group.instance.id]
5     user_data      = data.template_file.user_data.rendered
6
7     lifecycle {
8         create_before_destroy = true
9     }
10 }
```

Final Notes



Correct isolation, locking and state must be a priority

- Bugs in a program only break a part of an app
- Bugs in infrastructure can have catastrophic effects and result in whole systems crashing and becoming unworkable



Infrastructure has to be planned and incrementally tested

- We never code infrastructure "on the fly"



We never experiment with infrastructure in a production environment

- Always work in a sandbox
- With IaaS, this is easily done